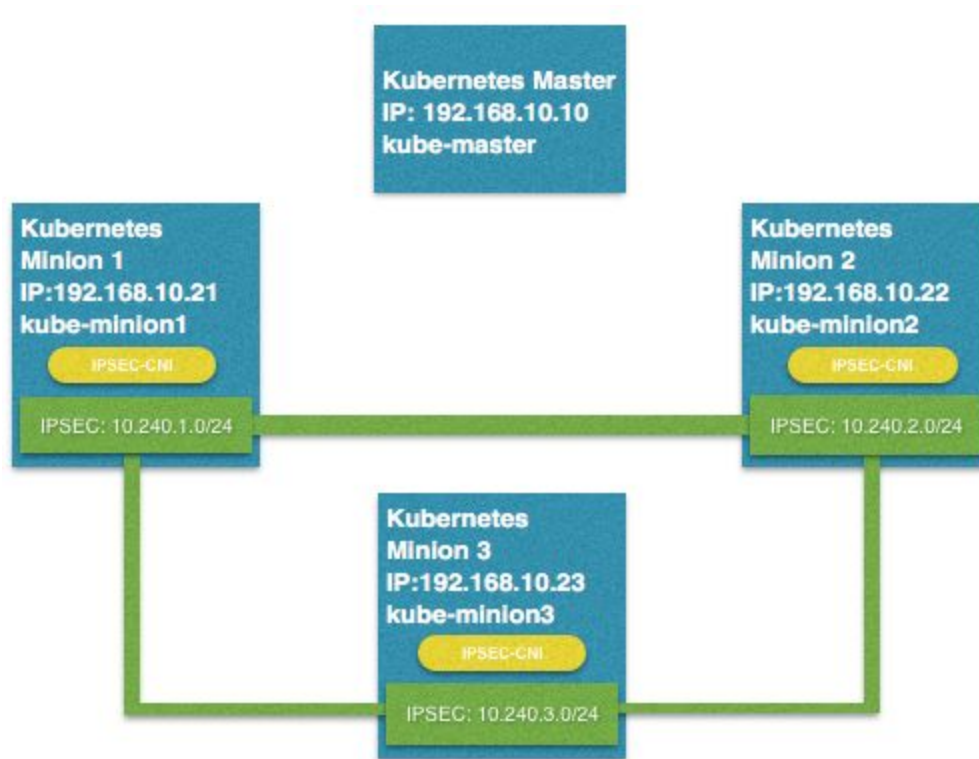


## Kubernetes CNI Project



The above diagram gives the basic layout of the setup.

Primary IPs

Master : 192.168.10.10

Minion1 : 192.168.10.21

Minion2 : 192.168.10.22

Minion3 : 192.168.10.23

Service Cluster range : 10.30.0.0/24

IP range for POD: 10.240.0.0/16

This range will be split as follows between the nodes via vagrant(shell script) and will be used as input for ip assignment by the plugin. Further version can work similar to flannel to keep that data in etcd and make the split configurable and dynamic.

Minion1 : 10.240.1.0/24

Minion2 : 10.240.2.0/24

Minion3 : 10.240.3.0/24

These same ranges will be used for StrongSwan IPSEC communication.

- The kubernetes master will have the following services running

- etcd
- api server
- scheduler
- controller manager
- Appropriate certs and token configuration will be required
- The kubernetes minion will have the following services running
  - docker
  - kubelet
    - cni configs will also be done here
  - kube-proxy
  - Strong swan
    - The server certs and client certs used will be common but we can use different certs for different servers.
    - Every node will have configuration for the other two nodes like a host-to-host configuration.

### Logical View of Node, POD and CNI Implementation



### Notes for CNI Implementation

- The CNI Plugin will get triggered with env params (PATH, Container name, Container interface name, namespace and command)
- The CNI plugin as standard input will get details such as bridge name(ipsec-br01) and ip cidr for pod per node (populated via Vagrant)
- The CNI plugin will check if the bridge exists, if not it will create and set first ip as the bridge ip.
- Further on it will create a veth pair (veth-truncated-CONTAINERNAME and container interface name given from env variables)
- It will place the veth-truncated-CONTAINERNAME in the ipse bridge and the other end it will move to container namespace

- The logic will further require to maintain a dictionary of container names ,container name spaces ,container interface names and ip addresses assigned. It will also store this data in the file.

## Sample Service Configurations

### StrongSwan IPSEC Minion1 (appropriate changes for Minion2 and 3)

# /etc/ipsec.conf - strongSwan IPsec configuration file for Minion1

```
config setup
conn %default
    ikelifetime=60m
    keylife=20m
    rekeymargin=3m
    keyingtries=1
    keyexchange=ikev2
    authby=psk

conn minion2
    left=%defaultroute
    leftsubnet=10.240.1.0/24
    leftid=minion1
    leftfirewall=yes
    rightid=minion2
    right=192.168.10.22
    rightsubnet=10.240.2.0/24
    auto=add

conn minion3
    left=%defaultroute
    leftsubnet=10.240.1.0/24
    leftid=minion3
    leftfirewall=yes
    rightid=minion3
    right=192.168.10.23
    rightsubnet=10.240.3.0/24
    auto=add
```

### ETCD

[Unit]

Description=etcd

Documentation=<https://github.com/coreos>

[Service]

```
ExecStart=/usr/bin/etcd --name $NODENAME \
    --cert-file=/etc/etcd/kubernetes.pem \
    --key-file=/etc/etcd/kubernetes-key.pem \
    --peer-cert-file=/etc/etcd/kubernetes.pem \
    --peer-key-file=/etc/etcd/kubernetes-key.pem \
    --trusted-ca-file=/etc/etcd/ca.pem \
    --peer-trusted-ca-file=/etc/etcd/ca.pem \
    --initial-advertise-peer-urls https://$IP:2380 \
    --listen-peer-urls https://$MASTER_IP:2380 \
    --listen-client-urls https://$MASTER_IP:2379,http://127.0.0.1:2379 \
    --advertise-client-urls https://$MASTER_IP:2379 \
    --initial-cluster-token etcd-cluster-0 \
    --initial-cluster $INITIAL_CLUSTER \
    --initial-cluster-state new \
    --data-dir=/var/lib/etcd
```

Restart=on-failure  
RestartSec=5

[Install]  
WantedBy=multi-user.target

### **API Server**

[Unit]  
Description=Kubernetes API Server

[Service]  
ExecStart=/usr/local/bin/kube-apiserver \  
--admission-control=NamespaceLifecycle,LimitRanger,SecurityContextDeny,ServiceAccount,ResourceQuota \  
--advertise-address=\${IP} \  
--allow-privileged=true \  
--apiserver-count=3 \  
--authorization-mode=ABAC \  
--authorization-policy-file=/var/lib/kubernetes/authorization-policy.json \  
--bind-address=0.0.0.0 \  
--enable-swagger-ui=true \  
--etcd-cafile=/var/lib/kubernetes/ca.pem \  
--insecure-bind-address=0.0.0.0 \  
--kubelet-certificate-authority=/var/lib/kubernetes/ca.pem \  
--etcd-servers=\${ETCD\_CLIENT\_ACCESS} \  
--service-account-key-file=/var/lib/kubernetes/kubernetes-key.pem \  
--service-cluster-ip-range=10.30.0.0/24 \  
--service-node-port-range=30000-32767 \  
--tls-cert-file=/var/lib/kubernetes/kubernetes.pem \  
--tls-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \  
--token-auth-file=/var/lib/kubernetes/token.csv \  
--v=2  
Restart=on-failure  
RestartSec=5

[Install]  
WantedBy=multi-user.target

### **Controller Manager**

[Unit]  
Description=Kubernetes Controller Manager

[Service]  
ExecStart=/usr/local/bin/kube-controller-manager \  
--allocate-node-cidrs=true \  
--cluster-name=kubernetes \  
--leader-elect=true \  
--master=http://\${MASTER\_IP}:8080 \  
--root-ca-file=/var/lib/kubernetes/ca.pem \  
--service-account-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \  
--service-cluster-ip-range=10.30.0.0/16 \  
--v=2  
Restart=on-failure  
RestartSec=5

[Install]  
WantedBy=multi-user.target

### **Scheduler**

[Unit]  
Description=Kubernetes Scheduler

```
[Service]
ExecStart=/usr/local/bin/kube-scheduler \
--leader-elect=true \
--master=http://$MASTER_IP:8080 \
--v=2
Restart=on-failure
RestartSec=5
```

```
[Install]
WantedBy=multi-user.target
```

## **Docker**

```
[Unit]
Description=Docker Application Container Engine
Documentation=http://docs.docker.io
```

```
[Service]
ExecStart=/usr/bin/docker daemon \
--iptables=false \
--ip-masq=false \
--host=unix:///var/run/docker.sock \
--log-level=error \
--storage-driver=overlay
Restart=on-failure
RestartSec=5
```

```
[Install]
WantedBy=multi-user.target
```

## **Kubelet**

```
[Unit]
Description=Kubernetes Kubelet
```

```
After=docker.service
Requires=docker.service
```

```
[Service]
ExecStart=/usr/bin/kubelet \
--allow-privileged=true \
--api-servers=${API_SERVERS} \
--cluster-dns=10.30.0.10 \
--cluster-domain=cluster.local \
--container-runtime=docker \
--network-plugin=ipsec-cni \
--kubeconfig=/var/lib/kubelet/kubeconfig \
--serialize-image-pulls=false \
--tls-cert-file=/var/lib/kubernetes/kubernetes.pem \
--tls-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \
--v=2
```

```
Restart=on-failure
RestartSec=5
```

```
[Install]
WantedBy=multi-user.target
```

## **kube-proxy**

```
[Unit]
Description=Kubernetes Kube Proxy
```

[Service]

ExecStart=/usr/bin/kube-proxy \

--masquerade-all \

--master=\$MASTER\_IP \

--kubeconfig=/var/lib/kubelet/kubeconfig \

--proxy-mode=iptables \

--v=2

Restart=on-failure

RestartSec=5

[Install]

WantedBy=multi-user.target